

numpy-intro

April 10, 2019

1 Numpy Intro

Felix Biessmann

2 numpy

- In machine learning you will work *a lot* with matrices
- numpy is the de facto standard for matrices in python
- Many other libraries have adopted and extended the API of numpy
 - scipy
 - pandas
 - many deep learning libraries

3 numpy is fast

```
In [174]: my_list = list(range(1000000))
          %timeit [x * 2 for x in my_list]
```

117 ms ± 18.8 ms per loop (mean ± std. dev. of 7 runs, 10 loops each)

```
In [2]: import numpy as np
        np.random.seed(0) # seed for reproducibility
        my_arr = np.arange(1000000)
        %timeit my_arr * 2
```

668 µs ± 79.2 µs per loop (mean ± std. dev. of 7 runs, 1000 loops each)

4 Creating Arrays

```
In [4]: np.array([1, 4, 2, 5, 3])
```

```
Out[4]: array([1, 4, 2, 5, 3])
```

		axis 1		
		0	1	2
axis 0	0	0,0	0,1	0,2
	1	1,0	1,1	1,2
	2	2,0	2,1	2,2

Array indexing

4.1 Specifying data types at array creation time

```
In [8]: np.array([1, 2, 3, 4], dtype='float32')
```

```
Out[8]: array([1., 2., 3., 4.], dtype=float32)
```

4.2 Creating two-dimensional arrays

```
In [14]: two_d_data = [list(range(i, i + 3)) for i in [2, 4, 6]]
          two_d_data
```

```
Out[14]: [[2, 3, 4], [4, 5, 6], [6, 7, 8]]
```

```
In [15]: np.array(two_d_data)
```

```
Out[15]: array([[2, 3, 4],
                [4, 5, 6],
                [6, 7, 8]])
```

4.3 Creating Arrays from Scratch

```
In [16]: np.zeros(10, dtype=int)
```

```
Out[16]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
In [28]: # Create a 3x5 floating-point array filled with ones
          np.ones((2, 3), dtype=float)
```

```
Out[28]: array([[1., 1., 1.],
                [1., 1., 1.]])
```

```
In [19]: # something like builtin range function
          np.arange(0, 20, 2)
```

```
Out[19]: array([ 0,  2,  4,  6,  8, 10, 12, 14, 16, 18])
```

```

In [20]: # Create an array of five values evenly spaced between 0 and 1
         np.linspace(0, 1, 5)

Out[20]: array([0.   , 0.25, 0.5  , 0.75, 1.   ])

In [39]: # Create a 3x3 array of values uniformly distributed between 0 and 1
         np.random.random((2, 3))

Out[39]: array([[0.5488135 , 0.71518937, 0.60276338],
                [0.54488318, 0.4236548 , 0.64589411]])

In [41]: # Create a 3x3 array of random integers in the interval [0, 10)
         np.random.randint(0, 10, (2, 3))

Out[41]: array([[8, 1, 5],
                [9, 8, 9]])

```

4.4 Array Indexing and Slicing

Arrays are indexed like this:

```
x[index]
```

and sliced like this:

```
x[start:stop:step]
```

With defaults `start=0`, `stop=size of dimension`, `step=1`.
Slices are *views* not *copies* (as for Lists)

4.4.1 Two Dimensional Arrays

Items can be accessed using a comma-separated tuple of indices:

```

In [54]: x2

Out[54]: array([[1, 5, 3, 9],
                [7, 2, 9, 8],
                [5, 0, 2, 8]])




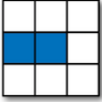
In [55]: x2[0, 0]

Out[55]: 1

In [56]: x2[2, -1]

Out[56]: 8

```

	Expression	Shape
	<code>arr[:2, 1:]</code>	<code>(2, 2)</code>
	<code>arr[2]</code> <code>arr[2, :]</code> <code>arr[2:, :]</code>	<code>(3,)</code> <code>(3,)</code> <code>(1, 3)</code>
	<code>arr[:, :2]</code>	<code>(3, 2)</code>
	<code>arr[1, :2]</code> <code>arr[1:2, :2]</code>	<code>(2,)</code> <code>(1, 2)</code>

Array indexing

4.4.2 Slicing Two Dimensional Arrays

Same as with lists and one-dimensional arrays:

```
In [57]: x2
```

```
Out[57]: array([[1, 5, 3, 9],
               [7, 2, 9, 8],
               [5, 0, 2, 8]])
```

```
In [58]: x2[0,:] # entire first row
```

```
Out[58]: array([1, 5, 3, 9])
```

```
In [59]: x2[0] # same as x2[0, :]
```

```
Out[59]: array([1, 5, 3, 9])
```

```
In [60]: x2[:,0] # entire first column
```

```
Out[60]: array([1, 7, 5])
```

4.5 Boolean Indexing

You can efficiently index arrays using boolean masks

```
In [116]: grid = np.arange(1, 10)
          grid
```

```
Out[116]: array([1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [115]: grid > 5
```

```
Out[115]: array([False, False, False, False, False,  True,  True,  True,  True])
```

```
In [117]: grid[grid > 5]
```

```
Out[117]: array([6, 7, 8, 9])
```

4.6 Fancy Indexing

Can give you arbitrary subsets of arrays given a set of indices

```
In [118]: x = np.random.randint(100, size=10)
          print(x)
```

```
[24 80 66 27 81 70 35 67 65 88]
```

```
In [119]: [x[1], x[4]]
```

```
Out[119]: [80, 81]
```

```
In [123]: ind = [1,4]
          x[ind]
```

```
Out[123]: array([80, 81])
```

5 Reshaping

Often you need to change the shape of an array; this is done with reshape

```
In [67]: grid = np.arange(1, 10)
          grid
```

```
Out[67]: array([1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [68]: grid = grid.reshape((3, 3))
          grid
```

```
Out[68]: array([[1, 2, 3],
                [4, 5, 6],
                [7, 8, 9]])
```

```
In [69]: row_vector = np.arange(3).reshape((1, 3))
          row_vector
```

```
Out[69]: array([[0, 1, 2]])
```

```
In [70]: column_vector = np.arange(3).reshape((3, 1))
          column_vector
```

```
Out[70]: array([[0],
                [1],
                [2]])
```

6 Broadcasting

```
In [5]: # elementwise addition
a = np.array([0, 1, 2])
a + 5
```

```
Out[5]: array([5, 6, 7])
```

```
In [6]: # elementwise addition
b = np.array([5, 5, 5])
a + b
```

```
Out[6]: array([5, 6, 7])
```

```
In [7]: M = np.ones((3, 3))
M
```

```
Out[7]: array([[1., 1., 1.],
               [1., 1., 1.],
               [1., 1., 1.]])
```

```
In [8]: M + a
```

```
Out[8]: array([[1., 2., 3.],
               [1., 2., 3.],
               [1., 2., 3.]])
```

```
In [10]: b = np.array([5, 5, 5]).reshape((3,1))
b
```

```
Out[10]: array([[5],
                [5],
                [5]])
```

```
In [13]: a + b
```

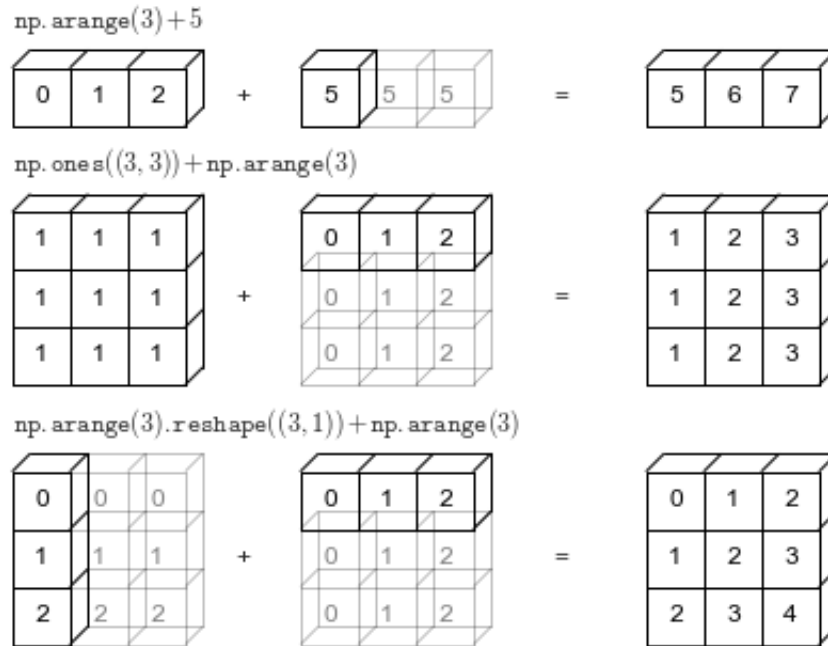
```
Out[13]: array([[5, 6, 7],
                [5, 6, 7],
                [5, 6, 7]])
```

```
In [15]: c = np.array([2,3])
a + c
```

ValueError

Traceback (most recent call last)

```
<ipython-input-15-ce668e4220e4> in <module>
1 c = np.array([2,3])
```



Broadcasting Visual

----> 2 a + c

ValueError: operands could not be broadcast together with shapes (3,) (2,)

In [17]: `a.reshape((3,1))`

Out[17]: `array([[0],
[1],
[2]])`

In [18]: `a.reshape((3,1)) + c`

Out[18]: `array([[2, 3],
[3, 4],
[4, 5]])`

7 Fast Vectorized Computations with UFuncs

- Python's default implementation CPython is slow for repeated executions
- This is mostly due to the fact that it's not compiled down to bytecode
- Major Bottlenecks are: type-checking and function dispatches
- numpy is fast since it allows to *vectorize* operations through NumPy's *universal functions* (ufuncs)
- For many tasks you can use plain python and numpy for efficient computations

7.1 ufunc Arithmetic Operators and Shortcuts

Operator	Equivalent ufunc	Description
+	np.add	Addition (e.g., $1 + 1 = 2$)
-	np.subtract	Subtraction (e.g., $3 - 2 = 1$)
-	np.negative	Unary negation (e.g., -2)
*	np.multiply	Multiplication (e.g., $2 * 3 = 6$)
/	np.divide	Division (e.g., $3 / 2 = 1.5$)
//	np.floor_divide	Floor division (e.g., $3 // 2 = 1$)
**	np.power	Exponentiation (e.g., $2 ** 3 = 8$)
%	np.mod	Modulus/remainder (e.g., $9 \% 4 = 1$)

8 Aggregation Functions

When working with large data sets, aggregations help to understand your data:

- Summing all values
- Min/Max
- Quantiles
- Mean/Median

Again, python itself is slow at that, but numpy is fast

```
In [86]: L = np.arange(10)
         L
```

```
Out[86]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
In [87]: sum(L)
```

```
Out[87]: 45
```

```
In [88]: np.sum(L)
```

```
Out[88]: 45
```

Function Name	NaN-safe Version	Description
np.sum	np.nansum	Compute sum of elements
np.prod	np.nanprod	Compute product of elements
np.mean	np.nanmean	Compute mean of elements
np.std	np.nanstd	Compute standard deviation
np.var	np.nanvar	Compute variance
np.min	np.nanmin	Find minimum value
np.max	np.nanmax	Find maximum value
np.argmin	np.nanargmin	Find index of minimum value
np.argmax	np.nanargmax	Find index of maximum value
np.median	np.nanmedian	Compute median of elements
np.percentile	np.nanpercentile	Compute rank-based statistics of elements
np.any	N/A	Evaluate whether any elements are true
np.all	N/A	Evaluate whether all elements are true

8.0.1 Multi dimensional aggregates

```
In [100]: M = np.ones((3, 4))  
M
```

```
Out[100]: array([[1., 1., 1., 1.],  
                [1., 1., 1., 1.],  
                [1., 1., 1., 1.]])
```

```
In [101]: M.sum()
```

```
Out[101]: 12.0
```

```
In [102]: M.sum(axis=0)
```

```
Out[102]: array([3., 3., 3., 3.])
```

```
In [103]: M.sum(axis=1)
```

```
Out[103]: array([4., 4., 4.])
```

8.1 Fast Sorting in NumPy: np.sort and np.argsort

```
In [124]: np.sort(heights)
```

```
Out[124]: array([163, 168, 168, 170, 170, 171, 173, 173, 173, 173, 174, 175, 175,  
                177, 178, 178, 178, 178, 179, 180, 182, 182, 182, 182, 183, 183,  
                183, 183, 183, 183, 183, 183, 185, 185, 185, 188, 188, 188, 189,  
                189, 193, 193])
```

```
In [125]: heights[heights.argsort()]
```

```
Out[125]: array([163, 168, 168, 170, 170, 171, 173, 173, 173, 173, 174, 175, 175,  
                177, 178, 178, 178, 178, 179, 180, 182, 182, 182, 182, 183, 183,  
                183, 183, 183, 183, 183, 183, 185, 185, 185, 188, 188, 188, 189,  
                189, 193, 193])
```

9 Saving and Loading Arrays

```
In [126]: arr = np.arange(10)  
          np.save('some_array', arr)  
          loaded_arr = np.load('some_array.npy')  
          loaded_arr
```

```
Out[126]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

9.1 Matrix Multiplication

```
In [167]: A = np.array([[1,1],[1,1]])
          B = np.array([[2,3],[4,5]])
          A.dot(B)
```

```
Out[167]: array([[6, 8],
                 [6, 8]])
```

```
In [168]: A @ B
```

```
Out[168]: array([[6, 8],
                 [6, 8]])
```

```
In [169]: (A @ B).T
```

```
Out[169]: array([[6, 6],
                 [8, 8]])
```

10 Visualization

Let's create some toy data:

$$X \in \mathbb{R}^{100,2} \sim \mathcal{N}(0,1)$$

```
In [4]: %matplotlib inline
import matplotlib.pyplot as plt
# set some other plotting style
plt.style.use('ggplot')

plt.figure(figsize=(6,6))
X = np.random.randn(100,2)
# add some offset
X += np.array([2,.5])
# scale the first dimension
X[:,0] *= 4
plt.scatter(X[:,0],X[:,1]);
plt.axis('equal');
```

