

# Lecture-5-Model-Evaluation

May 14, 2019

## 1 Machine Learning

Classification Metrics and Model Evaluation


## 2 Machine Learning Pipelines

## 3 Model Evaluation

- Machine Learning models are evaluated by comparing the predictions  $f(x) = \hat{y}$  and the target values  $y$
- Useful Metrics
  - Regression (for continuous predictions - covered later)
  - Classification
- As ML models can memorize any training data set, **all metrics must always be computed using cross-validation**

### 3.1 Classification Metrics

- Accuracy
- Precision
- Recall
- F1



ml-pipeline-2.png

## 3.2 Accuracy

Defined as  
$$\frac{\text{number of correct assignments}}{\text{number of data points}}$$
  
Problem: Imbalanced classes

```
In [53]: from sklearn.metrics import accuracy_score
```

```
    y_predicted = np.array([1, 0, 0, 0, 0])
    y_true = np.array([0, 1, 0, 0, 0])
```

```
In [54]: print(f"Accuracy: {accuracy_score(y_true, y_predicted)}")
```

Accuracy: 0.6

## 3.3 Precision

$$\text{precision} = \frac{|\{\text{relevant instances}\} \cap \{\text{predicted instances}\}|}{|\{\text{predicted instances}\}|}$$

```
In [55]: from sklearn.metrics import precision_score
```

```
    y_predicted = np.array([1, 0, 0, 0, 0])
    y_true = np.array([0, 1, 0, 0, 0])
```

```
In [56]: print(f"Accuracy: {accuracy_score(y_true, y_predicted)}")
        print(f"Precision: {precision_score(y_true, y_predicted)}")
```

Accuracy: 0.6

Precision: 0.0

```
In [57]: y_predicted = np.array([1, 1, 0, 0, 0])
    y_true = np.array([0, 1, 0, 0, 0])
```

```
In [58]: print(f"Accuracy: {accuracy_score(y_true, y_predicted)}")
        print(f"Precision: {precision_score(y_true, y_predicted)}")
```

Accuracy: 0.8

Precision: 0.5

## 3.4 Recall

$$\text{recall} = \frac{|\{\text{relevant instances}\} \cap \{\text{predicted instances}\}|}{|\{\text{relevant instances}\}|}$$

```
In [59]: from sklearn.metrics import recall_score
```

```
    y_predicted = np.array([1, 0, 0, 0, 0])
    y_true = np.array([0, 1, 0, 0, 0])
```

```
In [61]: print(f"Accuracy: {accuracy_score(y_true, y_predicted)}")
         print(f"Precision: {precision_score(y_true, y_predicted)}")
         print(f"Recall: {recall_score(y_true, y_predicted)}")
```

```
Accuracy: 0.6
Precision: 0.0
Recall: 0.0
```

```
In [62]: y_predicted = np.array([1, 1, 0, 0, 0])
         y_true = np.array([0, 1, 0, 0, 0])
```

```
In [63]: print(f"Accuracy: {accuracy_score(y_true, y_predicted)}")
         print(f"Precision: {precision_score(y_true, y_predicted)}")
         print(f"Recall: {recall_score(y_true, y_predicted)}")
```

```
Accuracy: 0.8
Precision: 0.5
Recall: 1.0
```

### 3.5 Precision and Recall

Source Wikipedia page on [Precision and Recall](#)

### 3.6 F1 Score

$$F = 2 \cdot \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$

```
In [71]: from sklearn.metrics import f1_score
```


```
         y_predicted = np.array([1, 0, 0, 0, 0])
         y_true = np.array([0, 1, 0, 0, 0])
```

```
In [72]: print(f"Precision: {precision_score(y_true, y_predicted)}")
         print(f"Recall: {recall_score(y_true, y_predicted)}")
         print(f"F1: {f1_score(y_true, y_predicted)}")
```

```
Precision: 0.0
Recall: 0.0
F1: 0.0
```

```
In [66]: y_predicted = np.array([1, 1, 0, 0, 0])
         y_true = np.array([0, 1, 0, 0, 0])
```

```
In [70]: print(f"Precision: {precision_score(y_true, y_predicted)}")
         print(f"Recall: {recall_score(y_true, y_predicted)}")
         print(f"F1: {f1_score(y_true, y_predicted)}")
```



```
Precision: 0.5
Recall: 1.0
F1: 0.6666666666666666
```

## 4 Model Evaluation

### 4.1 Cross-Validation

- ML models can memorize any data set (**overfitting**)
- But we want our models to perform well on new data (**generalization of learned rules**)
- Cross-validation emulates the setting of new unseen data:
  - Split data in training and test
  - Train model on training set
  - Test model on test set

### 4.2 Simplest Cross-Validation for Model Evaluation

Split data into **two folds** and train / test **just once**

---

Fold 1	Test
Fold 2	Train

---


### 4.3 Example 3-Fold Cross-Validation for Model Evaluation:

Split data into three folds:

— 4 —

---

Fold 1  
Fold 2



ml-pipeline-2.png

---

Fold 1	Train
Fold 2	Test
Fold 3	Train

---

#### 4.3.3 Fold 3 is Test Data

---

Fold 1	Train
Fold 2	Train
Fold 3	Test

---

## 5 Recap: Cross-Validation for Model Evaluation

- Split dataset in training and test set
- Requires **at least two** different folds / partitions / data sets (train and test)
- Fast Cross-validation:
  - Make two partitions, train and test
- Cross-validation with best generalization performance estimates:
  - k-fold

## 6 Example Application: Classifying Parliament Speeches

- Task: Political bias prediction from text data
- Data: [parliament speeches in German Bundestag](#)

## 7 Machine Learning Pipeline

```
In [3]: import os, gzip
import pandas as pd
import numpy as np
import urllib.request

import warnings
warnings.filterwarnings('ignore')

DATADIR = "data"

if not os.path.exists(DATADIR):
    os.mkdir(DATADIR)

file_name = os.path.join(DATADIR, 'bundestags_parlamentsprotokolle.csv.gzip')
if not os.path.exists(file_name):
    url_data = 'https://www.dropbox.com/s/1nlbfehnrrwa2zj/bundestags_parlamentsprotoko
    urllib.request.urlretrieve(url_data, file_name)

df = pd.read_csv(gzip.open(file_name), index_col=0).sample(frac=1)

alle_sprecher = df.sprecher.unique()
parteien = df.partei.unique()
partei_farben = {'cdcsu':'black', 'linke':'purple', 'spd':'red', 'gruene':'green', 'f

In [2]: df[:5]
```

Out [2]:

	sitzung	wahlperiode	sprecher \
30289	76	18	Dr.äKonstantin von Notz
11201	133	17	Peter Wichtel
24223	10	18	Dr.äEva Högl
34597	127	18	Dr. Joachim Pfeiffer
7715	96	17	Bettina Kudla

	text	partei
30289	Ich kann Ihnen schon sagen, wie das Ihre Unabh...	gruene
11201	Denn damit haben Sie am Ende genau das, was hi...	cdcsu
24223	Unser Koalitionsvertrag ist voller guter Ideen...	spd
34597	Frau Präsidentin! Liebe Kolleginnen und Kolleg...	cdcsu
7715	Zur Gläubigerbeteiligung - auch dieses Thema w...	cdcsu

```
In [31]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.neighbors import NearestCentroid
from sklearn.linear_model import SGDClassifier
from sklearn.pipeline import Pipeline
from sklearn.metrics import confusion_matrix, classification_report
```

```
In [22]: # Put some data aside for model evaluation
train_data, test_data, train_labels, test_labels = train_test_split(df['text'], df['p

ncc_clf = Pipeline([('vect', TfidfVectorizer(max_features=int(1e8))),
                    ('clf', NearestCentroid())]).fit(train_data, train_labels)

logreg_clf = Pipeline([('vect', TfidfVectorizer(max_features=int(1e8))),
                       ('clf', SGDClassifier())]).fit(train_data, train_labels)
```

## 7.1 Evaluating NCC on Training Data

```
In [23]: ncc_predictions = ncc_clf.predict(train_data)
         print(classification_report(ncc_predictions, train_labels))
```

	precision	recall	f1-score	support
cducsu	0.49	0.65	0.56	9612
fdp	0.42	0.21	0.29	5275
gruene	0.43	0.34	0.38	6323
linke	0.62	0.37	0.46	8222
spd	0.28	0.48	0.35	5511
micro avg	0.44	0.44	0.44	34943
macro avg	0.45	0.41	0.41	34943
weighted avg	0.47	0.44	0.43	34943

## 7.2 Evaluating NCC on Test Data

```
In [27]: ncc_predictions_test = ncc_clf.predict(test_data)
         print(classification_report(ncc_predictions_test, test_labels))
```

	precision	recall	f1-score	support
cducsu	0.48	0.61	0.54	2498
fdp	0.35	0.18	0.23	1365
gruene	0.40	0.32	0.35	1569
linke	0.54	0.34	0.42	1894
spd	0.26	0.45	0.33	1410
micro avg	0.41	0.41	0.41	8736
macro avg	0.41	0.38	0.38	8736
weighted avg	0.42	0.41	0.40	8736

### 7.3 Evaluating Logistic Regression on Training Data

```
In [26]: logreg_predictions = logreg_clf.predict(train_data)
         print(classification_report(logreg_predictions, train_labels))
```

	precision	recall	f1-score	support
cducsu	0.96	0.66	0.78	18852
fdp	0.27	0.98	0.42	738
gruene	0.69	0.93	0.79	3708
linke	0.85	0.87	0.86	4768
spd	0.62	0.85	0.72	6877
micro avg	0.76	0.76	0.76	34943
macro avg	0.68	0.86	0.71	34943
weighted avg	0.84	0.76	0.77	34943

### 7.4 Evaluating Logistic Regression on Training Data

```
In [28]: logreg_predictions_test = logreg_clf.predict(test_data)
         print(classification_report(logreg_predictions_test, test_labels))
```

	precision	recall	f1-score	support
cducsu	0.91	0.55	0.69	5273
fdp	0.04	0.62	0.08	50
gruene	0.33	0.62	0.43	658
linke	0.63	0.64	0.63	1177
spd	0.37	0.57	0.45	1578
micro avg	0.57	0.57	0.57	8736
macro avg	0.46	0.60	0.46	8736
weighted avg	0.72	0.57	0.61	8736

### 7.5 Summary Comparison NCC and Logistic Regression

Nearest Centroid Classifiers - NCC is a simple model - Overall performance not great - But not a lot of overfitting

Logistic Regression - More powerful model - Better prediction performance - More overfitting