

Machine Learning


Lecture 7 Perceptrons

Felix Bießmann

Beuth University & Einstein Center for Digital Future



Recap: Nearest Centroid Classification



Prototypes μ_{Δ} and μ_o can be the class means

$$\mu_{\Delta} = 1/N_{\Delta} \sum_n^{N_{\Delta}} \mathbf{x}_{\Delta,n}$$

$$\mu_o = 1/N_o \sum_n^{N_o} \mathbf{x}_{o,n}$$

Distance from w_{Δ} to new data \mathbf{x}

$$\|\mu_{\Delta} - \mathbf{x}\|_2$$



From Prototypes to Linear Classification

$$\begin{aligned} \text{distance}(\mathbf{x}, \mu_{\Delta}) &> \text{distance}(\mathbf{x}, \mu_o) \\ \|\mathbf{x} - \mu_{\Delta}\| &> \|\mathbf{x} - \mu_o\| \end{aligned} \tag{1}$$



From Prototypes to Linear Classification

$$\text{distance}(\mathbf{x}, \mu_{\Delta}) > \text{distance}(\mathbf{x}, \mu_o) \quad (1)$$

$$\|\mathbf{x} - \mu_{\Delta}\| > \|\mathbf{x} - \mu_o\|$$

$$\Leftrightarrow \|\mathbf{x} - \mu_{\Delta}\|^2 > \|\mathbf{x} - \mu_o\|^2$$

$$\Leftrightarrow \mathbf{x}^T \mathbf{x} - 2\mu_{\Delta}^T \mathbf{x} + \mu_{\Delta}^T \mu_{\Delta} > \mathbf{x}^T \mathbf{x} - 2\mu_o^T \mathbf{x} + \mu_o^T \mu_o$$

$$\Leftrightarrow \mu_{\Delta}^T \mathbf{x} - \mu_{\Delta}^2/2 < \mu_o^T \mathbf{x} - \mu_o^2/2$$

$$\Leftrightarrow 0 < \underbrace{(\mu_o - \mu_{\Delta})^T}_{\mathbf{w}} \mathbf{x} - 1/2 \underbrace{(\mu_o^T \mu_o - \mu_{\Delta}^T \mu_{\Delta})}_{\beta}$$



From Prototypes to Linear Classification

$$\text{distance}(\mathbf{x}, \mu_{\Delta}) > \text{distance}(\mathbf{x}, \mu_o) \quad (1)$$

$$\|\mathbf{x} - \mu_{\Delta}\| > \|\mathbf{x} - \mu_o\|$$

$$\Leftrightarrow \|\mathbf{x} - \mu_{\Delta}\|^2 > \|\mathbf{x} - \mu_o\|^2$$

$$\Leftrightarrow \mathbf{x}^T \mathbf{x} - 2\mu_{\Delta}^T \mathbf{x} + \mu_{\Delta}^T \mu_{\Delta} > \mathbf{x}^T \mathbf{x} - 2\mu_o^T \mathbf{x} + \mu_o^T \mu_o$$

$$\Leftrightarrow \mu_{\Delta}^T \mathbf{x} - \mu_{\Delta}^2/2 < \mu_o^T \mathbf{x} - \mu_o^2/2$$


$$\Leftrightarrow 0 < \underbrace{(\mu_o - \mu_{\Delta})^T \mathbf{x}}_w - 1/2 \underbrace{(\mu_o^T \mu_o - \mu_{\Delta}^T \mu_{\Delta})}_{\beta}$$

Linear Classification

$$\mathbf{w}^T \mathbf{x} - \beta = \begin{cases} > 0 & \text{if } \mathbf{x} \text{ belongs to class } o \\ < 0 & \text{if } \mathbf{x} \text{ belongs to class } \Delta \end{cases} \quad (2)$$




Linear Classification



$$\mathbf{w}^T \mathbf{x} - \beta = \begin{cases} > 0 & \text{if } \mathbf{x} \text{ belongs to } o \\ < 0 & \text{if } \mathbf{x} \text{ belongs to } \Delta \end{cases}$$



Linear Classification



What is a good w ?

→ We need an **error function** that tells us how good w is.



Two classical Error Functions

Given data $\mathbf{x} \in \mathbb{R}^D$ and corresponding labels $y \in \{-1, +1\}$, two classical error functions $\mathcal{E}(\mathbf{x}, y, \mathbf{w})$ to find the optimal $\mathbf{w} \in \mathbb{R}^D$ are:

Error Function	Used in
$\frac{1}{2}(y - \mathbf{w}^\top \mathbf{x})^2$	Adaline [?]
$\max(0, -y\mathbf{w}^\top \mathbf{x})$	Perceptron [?]



Rosenblatt's Perceptron



Frank Rosenblatt
(1928-1969)

Rosenblatt proposed the **perceptron**, an artificial neural network for pattern recognition [?]

Perceptrons gave rise to the field of artificial neural networks




Rosenblatt's Perceptron

Rosenblatt studied Perceptrons, so that

“[...] the fundamental laws of organization which are common to all information handling systems, machines and men included, may eventually be understood.”



Artificial Neural Networks



Input nodes x_i receive information

Inputs are multiplied with a weighting factor w_i and summed up

Integrated input is mapped through some (non-linear) function $f(\cdot)$

$$f(x) = \begin{cases} +1 & \text{if } \mathbf{x} \text{ is preferred stimulus} \\ -1 & \text{if } \mathbf{x} \text{ is any other stimulus} \end{cases}$$



Rosenblatt's Perceptron



Perceptron Hardware (pictures from [?])



The Perceptron Learning Algorithm

Goal Binary classification of multivariate data $\mathbf{x} \in \mathbb{R}^D$



The Perceptron Learning Algorithm

Goal Binary classification of multivariate data $\mathbf{x} \in \mathbb{R}^D$

Input Learning rate η and N tuples (\mathbf{x}_n, y_n) where
 $\mathbf{x}_n \in \mathbb{R}^D$ is the D -dimensional data
 $y_n \in \{-1, +1\}$ is the corresponding label, such that

$$y_n = \begin{cases} +1 & \text{if } \mathbf{x}_n \text{ is a correct example} \\ -1 & \text{if } \mathbf{x}_n \text{ is **not** a correct example} \end{cases}$$



The Perceptron Learning Algorithm

Goal Binary classification of multivariate data $\mathbf{x} \in \mathbb{R}^D$

Input Learning rate η and N tuples (\mathbf{x}_n, y_n) where

$\mathbf{x}_n \in \mathbb{R}^D$ is the D -dimensional data

$y_n \in \{-1, +1\}$ is the corresponding label, such that

$$y_n = \begin{cases} +1 & \text{if } \mathbf{x}_n \text{ is a correct example} \\ -1 & \text{if } \mathbf{x}_n \text{ is \textbf{not} a correct example} \end{cases}$$

Output Weight vector $w \in \mathbb{R}^D$ such that

$$\mathbf{w}^\top \mathbf{x}_n = \begin{cases} \geq 0 & \text{if } y_n = +1 \\ < 0 & \text{if } y_n = -1 \end{cases}$$




The Perceptron Error Function

$$\mathcal{E}_{\mathcal{P}}(\mathbf{w}) = - \sum_{m \in \mathcal{M}} \mathbf{w}^{\top} \mathbf{x}_m y_m \quad (3)$$



Classification Error as Function of Weights





Given data $\mathbf{x} \in \mathbb{R}^D$ and corresponding labels $y \in \{-1, +1\}$ the classification error \mathcal{E} is a function of the weights \mathbf{w} (and the data \mathbf{x} , y)

$$\mathcal{E}(\mathbf{w}, \mathbf{x}_m, y_m) = - \sum_{m \in \mathcal{M}} \mathbf{w}^\top \mathbf{x}_m y_m \quad (4)$$


where \mathcal{M} denotes the index set of all *misclassified* data \mathbf{x}_m



Classification Error as Function of Weights



Gradient Descent




How to minimize the error function?

$$\mathcal{E}(\mathbf{w}, \mathbf{x}_m, y_m) = - \sum_{m \in \mathcal{M}} \mathbf{w}^\top \mathbf{x}_m y_m$$

→ **Gradient Descent**



Gradient Descent



We minimize $\mathcal{E}(\mathbf{w}, \mathbf{x}_m, y_m)$ by walking in the opposite direction of the gradient.

$$\mathbf{w}^{\text{new}} \leftarrow \mathbf{w}^{\text{old}} - \eta \frac{1}{|\mathcal{X}|} \sum_{i=1}^{|\mathcal{X}|} \nabla \mathcal{E}(\mathbf{w}, \mathbf{x}_i, y_i)$$

where \mathcal{X} is the set of data points and η is called a **learning rate**.



Stochastic Gradient Descent

A noisy estimate of


$$\frac{1}{|\mathcal{X}|} \sum_{i=1}^{|\mathcal{X}|} \nabla \mathcal{E}(\mathbf{w}, \mathbf{x}_i, y_i)$$

is obtained by [?]

$$\mathbf{w}^{\text{new}} \leftarrow \mathbf{w}^{\text{old}} - \eta \nabla \mathcal{E}(\mathbf{w}, \mathbf{x}_i, y_i)$$

Note that only \mathbf{w} is stored and only one data point \mathbf{x}_i and label y_i are considered at a time!

→ Scales to large data sets [?]



Perceptron Training

Training a Perceptron means finding weights \mathbf{w} that minimize $\mathcal{E}_{\mathcal{P}}$:

$$\operatorname{argmin}_w \left(- \sum_{m \in \mathcal{M}} \mathbf{w}^\top \mathbf{x}_m y_m \right) \quad (5)$$

where \mathcal{M} denotes the index set of all *misclassified* data \mathbf{x}_m




Perceptron Training

Training a Perceptron means finding weights \mathbf{w} that minimize $\mathcal{E}_{\mathcal{P}}$:

$$\operatorname{argmin}_{\mathbf{w}} \left(- \sum_{m \in \mathcal{M}} \mathbf{w}^{\top} \mathbf{x}_m y_m \right) \quad (5)$$

where \mathcal{M} denotes the index set of all *misclassified* data \mathbf{x}_m
Data $\mathbf{x} \in \mathbb{R}^2$





Perceptron Training

Training a Perceptron means finding weights \mathbf{w} that minimize $\mathcal{E}_{\mathcal{P}}$:

$$\operatorname{argmin}_{\mathbf{w}} \left(- \sum_{m \in \mathcal{M}} \mathbf{w}^{\top} \mathbf{x}_m y_m \right) \quad (5)$$

where \mathcal{M} denotes the index set of all *misclassified* data \mathbf{x}_m



The Perceptron Learning Algorithm

Eq. 5 can be minimized *iteratively*
using **stochastic gradient descent**
[??]

1. Initialize \mathbf{w}^{old} (randomly, $1/n$, ...)
2. While there are misclassified data points \mathbf{x}_m
Pick a random misclassified data point \mathbf{x}_m

$$\mathbf{w}^{\text{new}} \leftarrow \mathbf{w}^{\text{old}} - \eta \nabla \mathcal{E}_{\mathcal{P}}(\mathbf{w}) = \mathbf{w}^{\text{old}} + \eta \mathbf{x}_m y_m \quad (6)$$



The Perceptron Learning Algorithm

Algorithm 1 Perceptron learning


Require: Data $X = [\mathbf{x}_1, \dots, \mathbf{x}_N]$, $\mathbf{x}_i \in \mathbb{R}^D$, Labels $y_1, \dots, y_N \in \{-1, +1\}$, iterations its , learning rate η

Ensure: \mathbf{w}


```
1:  $\mathbf{w} = \mathbf{1}_D / (D)$ 
2: for  $it = 1, \dots, its$  do
3:   # Pick a random example
4:    $i = \text{ceil}(\text{rand} * N)$ 
5:   # If this example is not correctly classified
6:   if  $\text{sign}(\mathbf{w}^\top \mathbf{x}_i) \neq y_i$  then
7:     # Update weight vector
8:      $\mathbf{w} \leftarrow \mathbf{w} + \eta / it \mathbf{x}_i y_i$ 
9:   end if
10: end for
```




The Perceptron Learning Algorithm in Action



The Perceptron Learning Algorithm in Action



The Perceptron Learning Algorithm in Action



Problems with Perceptrons



- Each update might lead to new misclassifications
- Many solutions might exist; which one is correct?
- Convergence might be slow
- Only solves linearly separable problems
- If there is no solution, the algorithm will not converge

Some solutions have been proposed by ?



Problems with Perceptrons

Perceptrons can only learn linearly separable problems.



Application example: Handwritten Digit Recognition

Handwritten digits
from UPS data set



Each digit represented as
 16×16 pixel image

→ $x \in \mathbb{R}^{256}$ input nodes

Each image is associated
with a label

$y \in \{0, 1, \dots, 9\}$


Goal Artificial neural network that
recognizes the digit 8

⇒ We need a function $f(\cdot)$
such that


$$f(x) = \begin{cases} -1 & \text{if } y \in \{0, 1, \dots, 7, 9\} \\ +1 & \text{if } y = 8 \end{cases}$$



Application example: Handwritten Digit Recognition



weight vector



Application example: Handwritten Digit Recognition

- Handwritten Digit Recognition is a **Multiclass Problem**
 - But the Perceptron is a **two-class** (binary) classifier
 - How can we **binarize** the multi class problem?
- Train a perceptron for each digit against all others
- Concatenate all weight vectors $W = [\mathbf{w}_0, \mathbf{w}_1, \dots, \mathbf{w}_9]$
- For new data point \mathbf{x} compute the label as $\operatorname{argmax}(W^T \mathbf{x})$



Minimizing Functions

Remember:

- ML algorithms are functions f that need to be fitted to data
- ML is finding minima/maxima of functions
- Function minimization is done by **Gradient Descent**



Taking Derivatives of Univariate Functions

Function	Derivative
<hr/>	
$f(x) = x^n$	

$$cf(x)$$

$$f(x) + g(x)$$

$$f(x)g(x)$$

$$\frac{f(x)}{g(x)}$$

$$f(g(x))$$



Taking Derivatives of Univariate Functions

Function	Derivative
$f(x) = x^n$	$f'(x) = nx^{n-1}$
$cf(x)$	$cf'(x)$
$f(x) + g(x)$	$f'(x) + g'(x)$
$f(x)g(x)$	$f'(x)g(x) + f(x)g'(x)$
$\frac{f(x)}{g(x)}$	$\frac{f'(x)g(x) - f(x)g'(x)}{g(x)^2}$
$f(g(x))$	$f'(g(x)) g'(x)$



Taking Derivatives of Vector-Valued Functions

Consider a vector-valued function $f(\mathbf{x}) : \mathbb{R}^D \rightarrow \mathbb{R}^U, \mathbf{x} \in \mathbb{R}^D$

The gradient $\nabla f(\mathbf{x})$ at position \mathbf{x} is

$$\nabla f(\mathbf{x}) = \left[\frac{\partial f(\mathbf{x})}{\partial x_1}, \dots, \frac{\partial f(\mathbf{x})}{\partial x_D} \right] \quad (7)$$



Taking Derivatives of Vector-Valued Functions

Function	Derivative
$f(\mathbf{x}) = \mathbf{u}^\top \mathbf{x}$	
$f(\mathbf{x}) = \mathbf{x}^\top \mathbf{x}$	
$f(\mathbf{x}) = A\mathbf{x}$	
$f(\mathbf{x}) = \mathbf{x}^\top A\mathbf{x}$	

More complicated examples: The Matrix Cookbook



Taking Derivatives of Vector-Valued Functions

Function	Derivative
$f(\mathbf{x}) = \mathbf{u}^\top \mathbf{x}$	$f'(\mathbf{x}) = \mathbf{u}$
$f(\mathbf{x}) = \mathbf{x}^\top \mathbf{x}$	$f'(\mathbf{x}) = 2\mathbf{x}$
$f(\mathbf{x}) = A\mathbf{x}$	$f'(\mathbf{x}) = A$
$f(\mathbf{x}) = \mathbf{x}^\top A\mathbf{x}$	$f'(\mathbf{x}) = \mathbf{x}^\top (A + A^\top)$

More complicated examples: The Matrix Cookbook



References

- C. M. Bishop. *Pattern Recognition and Machine Learning (Information Science and Statistics)*. Springer US, 2007.
- L. Bottou. Large-scale machine learning with stochastic gradient descent. In Y. Lechevallier and G. Saporta, editors, *Proceedings of the 19th International Conference on Computational Statistics (COMPSTAT'2010)*, pages 177–187, Paris, France, 2010. Springer.
- Y. Freund and R. E. Schapire. Large margin classification using the perceptron algorithm. *Machine Learning*, pages 277–296, 1998.
- H. Robbins and S. Monro. A stochastic approximation method. *Annals of Mathematical Statistics*, 22(3):400–407, 1951.
- F. Rosenblatt. The perceptron: a probabilistic model for information storage and organization in the brain. *Psychological Review*, 65(6):386–408, Nov. 1958.
- B. Widrow and M. E. Hoff. Adaptive switching circuits. In *1960 IRE WESCON Convention Record, Part 4*, pages 96–104, New York, 1960. IRE.

