

Machine Learning

Neural Networks


Felix Bießmann

Beuth University & Einstein Center for Digital Future

June 25, 2019



# Linear Classification





$$\phi(\mathbf{w}^T \mathbf{x} - \beta) = \begin{cases} > 0 & \text{if } \mathbf{x} \text{ is from class } o \\ < 0 & \text{if } \mathbf{x} \text{ is from class } \Delta \end{cases}$$




# Problems with Perceptrons

Perceptrons can only learn linearly separable problems.




# Deep Neural Networks



# Deep Neural Networks

Combinations of Perceptrons (Multi Layer Perceptrons):

Hidden Units



Input Units

Output Units

Neurons (Units), that are neither output nor input are called  
**Hidden Units.**



# A Short History of Deep Learning

- 1943 First mathematical Neuron Model (Mcculloch and Pitts, 1943)
- 1957 Perceptron Algorithm (Rosenblatt, 1958)
- 1969 Perceptrons cannot solve non-linearen Problems (Minsky and Papert, 1969)
- 1970 Backpropagation: Efficient gradient computations (Linnainmaa, 1970)
- 1980 Computer Hardware  $\approx 10,000$  faster compared to 1960/1970 – Automatic Differentiation (Speelpenning, 1980)
- 1986 Backpropagation learns meaningful representations (Rumelhart et al., 1986), NETtalk (Sejnowski and Rosenberg, 1986)
- 1992 Support-Vector Machines (SVMs) (Boser et al., 1992)
- 2000 Computer Hardware (GPUs)  $\approx 10,000$  faster compared to 1980/1990, Bigger datasets render kernel SVMs computationally infeasible
- 2012 Deep Convolutional Networks wins ImageNet (Krizhevsky et al., 2012)
- 2014 Neural Machine Translation surpasses traditional methods
- 2017 Neural Networks for Reinforcement Learning excell at Go (AlphaGo Zero)
- 2018 ImageNet Moment for Neural Language Models (BERT / ELMO)

Sources: Juergen Schmidhuber's page and others



# Universal Approximation Theorem

[Cybenko, 1989]

Multilayer Perceptrons with one hidden layer and a finite number of hidden units can approximate any function.



# Training of Deep Neural Networks

- Training: Gradient Descent
  - Problem: Gradient Computations
    - Mathematically challenging for complex models
    - Computationally challenging
- Solution: **Backpropagation**
- Elegant formulation
  - Efficient implementation





# Backpropagation Algorithm

---

## Algorithm 1 Backpropagation Algorithm - Pseudocode

---

**Require:** Data  $\mathbf{X} \in \mathbb{R}^{D \times N}$ , labels  $\mathbf{Y} \in \mathbb{R}^{K \times N}$ , untrained network


**Ensure:** Network parameters  $\mathbf{W}^{(1)}, \dots, \mathbf{W}^{(V)}$

```
1: while Not converged do  
2:   # Compute network predictions  
3:   # Evaluate error function  
4:   # Propagate error from output layer back to input layer  
5:   # Take gradient descent step  
6: end while
```

---




# Learning with Backpropagation in Neural Networks



Computation of network predictions is called **Forward Propagation**.




# Learning with Backpropagation in Neural Networks



**Backpropagation** refers to efficient computation of error function gradients for all connections.




# Learning with Backpropagation in Neural Networks



After a forward and backward pass a gradient step is performed.



# Forward Pass



Each neuron computes a weighted sum  $a_j$  of its inputs


$$a_j = \sum_i w_{ji}^{(v)} z_i \quad (1)$$

and transforms  $a_j$  with some non-linear function  $\phi(\cdot)$

$$z_j = \phi(a_j). \quad (2)$$



# Forward Pass



Each neuron computes a weighted sum  $a_j$  of its inputs

$$a_j = \sum_i w_{ji}^{(\nu)} z_i \quad (1)$$

and transforms  $a_j$  with some non-linear function  $\phi(\cdot)$


$$z_j = \phi(a_j). \quad (2)$$

Input Layer:  $z_i \equiv x_i$

Output Layer:  $z_i \equiv \hat{y}_i$



# Forward Pass



Each neuron computes a weighted sum  $a_j$  of its inputs

$$a_j = \sum_i w_{ji}^{(v)} z_i \quad (1)$$

and transforms  $a_j$  with some non-linear function  $\phi(\cdot)$


$$z_j = \phi(a_j). \quad (2)$$

After a forward pass the error function is evaluated:

$$J(\hat{\mathbf{y}}, \mathbf{y}) = \frac{1}{2} (\hat{\mathbf{y}} - \mathbf{y})^2 \quad (3)$$



# Error Backpropagation



**Goal:**


Computation of gradient of error function  $J$

$$\frac{\partial J(\mathbf{y}, \mathbf{x}, \mathbf{W}^{(1)}, \mathbf{W}^{(2)})}{\partial w_{ji}^{(v)}} \quad (4)$$





# Error Backpropagation



## Backpropagation Idea

$w_{ji}^{(\nu)}$  changes  $J$   
**only** through summed up inputs  $a_j$

Gradient of error function (chain rule):

$$\frac{\partial J}{\partial w_{ji}^{(\nu)}} = \frac{\partial J}{\partial a_j} \frac{\partial a_j}{\partial w_{ji}^{(\nu)}}$$



# Error Backpropagation


Error function gradient:

$$\frac{\partial J}{\partial w_{ji}^{(v)}} = \frac{\partial J}{\partial a_j} \frac{\partial a_j}{\partial w_{ji}^{(v)}} \quad (4)$$

From 1:

$$a_j = \sum_i w_{ji}^{(v)} z_i$$

$$\frac{\partial a_j}{\partial w_{ji}^{(v)}} = \frac{\partial \sum_i w_{ji}^{(v)} z_i}{\partial w_{ji}^{(v)}} \quad (5)$$
$$= z_i$$



Error Backpropagation



# Error Backpropagation

Error function gradient:


$$\frac{\partial J}{\partial w_{ji}^{(v)}} = \frac{\partial J}{\partial a_j} \underbrace{\frac{\partial a_j}{\partial w_{ji}}}_{z_i}$$

at output units ( $\mathbf{z} \equiv \mathbf{a} \equiv \hat{\mathbf{y}}$ ):


Outer derivative of  $J$ , e.g.

$$J = \frac{1}{2}(\hat{\mathbf{y}} - \mathbf{y})^2 \quad (4)$$

$$\frac{\partial J}{\partial a_j} = (\hat{\mathbf{y}} - \mathbf{y}) \equiv \delta_j \quad (5)$$



# Error Backpropagation



Error function gradient:

$$\frac{\partial J}{\partial w_{ji}^{(v)}} = \underbrace{\frac{\partial J}{\partial a_j}}_{\delta_j} \underbrace{\frac{\partial a_j}{\partial w_{ji}}}_{z_i}$$

$\delta_j$  is the error signal of the **receiving** neurons  $j$

$z_i$  is the activation of the **sending** neuron  $i$



# Error Backpropagation


Error function gradient:

$$\frac{\partial J}{\partial w_{ji}^{(v)}} = \underbrace{\frac{\partial J}{\partial a_j}}_{\delta_j} \underbrace{\frac{\partial a_j}{\partial w_{ji}}}_{z_i}$$

at hidden units:

## Backpropagation Idea

$a_j$  changes  $J$  **only** via outputs to  $a_k$



$$\delta_j \equiv \frac{\partial J}{\partial a_j} = \sum_k \frac{\partial J}{\partial a_k} \frac{\partial a_k}{\partial a_j} \quad (4)$$



# Error Backpropagation


Error function gradient:

$$\frac{\partial J}{\partial w_{ji}^{(v)}} = \underbrace{\frac{\partial J}{\partial a_j}}_{\delta_j} \underbrace{\frac{\partial a_j}{\partial w_{ji}}}_{z_i}$$


$\delta_j$  at hidden units:

$$\delta_j \equiv \frac{\partial J}{\partial a_j} = \sum_k \frac{\partial J}{\partial a_k} \frac{\partial a_k}{\partial a_j} \quad (4)$$

$$\begin{aligned} \frac{\partial a_k}{\partial a_j} &= \frac{\partial w_{kj}^{(v)} \phi(a_j)}{\partial a_j} \\ &= w_{kj}^{(v)} \phi'(a_j) \end{aligned} \quad (5)$$



# Error Backpropagation



Error function gradient:

$$\frac{\partial J}{\partial w_{ji}^{(v)}} = \underbrace{\frac{\partial J}{\partial a_j}}_{\delta_j} \underbrace{\frac{\partial a_j}{\partial w_{ji}}}_{z_i}$$

$\delta_j$  at hidden units:


$$\delta_j \equiv \frac{\partial J}{\partial a_j} = \sum_k \frac{\partial J}{\partial a_k} \frac{\partial a_k}{\partial a_j} \quad (4)$$

$$\frac{\partial a_k}{\partial a_j} = w_{kj}^{(v)} \phi'(a_j)$$

$$\frac{\partial J}{\partial a_k} = \delta_k$$



# Error Backpropagation



Error function gradient:

$$\frac{\partial J}{\partial w_{ji}^{(v)}} = \underbrace{\frac{\partial J}{\partial a_j}}_{\delta_j} \underbrace{\frac{\partial a_j}{\partial w_{ji}}}_{z_i}$$


$\delta_j$  at hidden units:

$$\delta_j \equiv \frac{\partial J}{\partial a_j} = \phi'(a_j) \sum_k w_{kj}^{(v)} \delta_k \quad (4)$$





# Error Backpropagation - Simple Example



Network with one hidden layer and linear output function

$$\hat{\mathbf{y}} = \mathbf{a}$$

$\phi(z)$  is the logistic function


$$\phi(z) = \frac{1}{1 + e^{-z}}$$

$J$  is the quadratic error

$$J = \frac{1}{2}(\hat{\mathbf{y}} - \mathbf{y})^2$$



# Error Backpropagation - Simple Example



For each datapoint  $\mathbf{x}_i$  the prediction  $\hat{\mathbf{y}}$  is computed


$$a_j = \sum_{i=0}^D w_{ji}^{(1)} x_i$$

$$z_j = \phi(a_j)$$

$$\hat{y}_k = \sum_{j=0}^M w_{kj}^{(2)} z_j$$



# Error Backpropagation - Simple Example



The error signal  $\delta$  is:  
at the output layer


$$\delta_k = \hat{y}_k - y_k$$

at hidden units

$$\begin{aligned}\delta_j &= \phi'(z_j) \sum_{k=1}^K w_{kj}^{(2)} \delta_k \\ &= \phi(z_j)(1 - \phi(z_j)) \sum_{k=1}^K w_{kj}^{(2)} \delta_k\end{aligned}$$



# Error Backpropagation - Simple Example



Full gradient

$$\frac{\partial J}{\partial w_{kj}^{(2)}} = \delta_k z_j$$

$$\frac{\partial J}{\partial w_{ji}^{(1)}} = \delta_j x_i$$



# Other Loss Functions

Error Function

Used in

---

$$\frac{1}{2}(y - \mathbf{w}^\top \mathbf{x})^2$$

Adaline [Widrow and Hoff, 1960]

$$\max(0, -y\mathbf{w}^\top \mathbf{x})$$

Perceptron [Rosenblatt, 1958]

$$-\sum_{k=1}^K y_{\text{true}} \log(y_{\text{predicted}})$$

Most classification neural networks



# Cross-Entropy

$$-\sum_{k=1}^K y_{\text{true}} \log(y_{\text{predicted}}) \quad (5)$$

Where

- $K$  is the number of classes
- $y_{\text{true}}$  is the one-hot encoded label
- $\mathbf{z}_k$  is the activity of the  $k$ th neuron in the last layer and

$$y_{\text{predicted}} = \frac{e^{\mathbf{z}_k}}{\sum_{k=1}^K e^{\mathbf{z}_k}} \quad (6)$$



# Backpropagation Algorithm

---

## Algorithm 2 Backpropagation Algorithm

---

**Require:** Data  $\mathbf{X} \in \mathbb{R}^{D \times N}$ , labels  $\mathbf{Y} \in \mathbb{R}^{K \times N}$ , untrained network

**Ensure:** network parameters  $\mathbf{W}^{(1)}, \dots, \mathbf{W}^{(V)}$

```
1: while Not converged do
2:   # Forward Propagation
3:   # Input Layer:
4:    $\mathbf{z}_0 = \phi(\mathbf{W}^{(0)}\mathbf{x}_i)$ 
5:   for Layer  $v = 1, \dots, V$  do
6:      $\mathbf{z}_v = \phi(\mathbf{W}^{(v)}\mathbf{z}_{v-1})$ 
7:   end for
8:   # Error Computation at Output Layer (quadratic error)
9:    $\delta_{V+1} = \mathbf{z}_V - \mathbf{y}_i$ 
10:  # Backpropagation
11:  for Layer  $v = V, \dots, 1$  do
12:    # Error Signal in Layer  $v$ 
13:     $\delta_v = \phi'(\mathbf{z}_v)^\top \delta_{v+1}^\top \mathbf{W}^{(v)}$ 
14:    # Gradient Step
15:     $\mathbf{W}^{(v)} = \mathbf{W}^{(v)} - \eta \delta_v \mathbf{z}_{v-1}^\top$ 
16:  end for
17: end while
```

---



## Summary

- Perceptrons cannot separate linearly non-separable problems
- Using combinations and stacking of standard Perceptrons, Multi Layer Perceptrons (MLPs) can approximate any function with one hidden layer
- Gradient descent for MLPs is challenging, mathematically and computationally
- Backpropagation: Efficient Gradient computation





# References

- C. M. Bishop. *Pattern Recognition and Machine Learning (Information Science and Statistics)*. Springer US, 2007.
- G. Cybenko. Approximations by superpositions of sigmoidal functions. *Mathematics of Control, Signals and Systems*, 2(4):303–314, 1989.
- T. Hastie, R. Tibshirani, and J. H. Friedman. *The Elements of Statistical Learning*. 2003.
- K. P. Murphy. *Machine Learning: A Probabilistic Perspective*. Adaptive Computation and Machine Learning. The MIT Press, 1 edition, 2012. ISBN 0262018020,9780262018029.
- F. Rosenblatt. The perceptron: a probabilistic model for information storage and organization in the brain. *Psychological Review*, 65(6):386–408, Nov. 1958.
- B. Widrow and M. E. Hoff. Adaptive switching circuits. In *1960 IRE WESCON Convention Record, Part 4*, pages 96–104, New York, 1960. IRE.

