

Workshop 3

Non-Linear Regression Models II

Start R in the usual way.

In the 2nd section you will be using functions from the package `gam` and `akima` which you will probably need to install. You will also use the package `ISLR`, but if you are up-to-date with the workshops, this will already be installed. The other packages used `MASS` and `splines` were downloaded when R was installed.

Download there two script files `loess1.r` and `loess2.r` from Moodle, which contain template code for Section 1.

1 Local regression

The code in first script file `loess1.r` uses local regression to estimate the linear loess value at a specific point x_0 . The code in `loess2.r` does the same over a given grid of x -values

Last week you used non-linear smoothing on the motorcycle helmet acceleration data `mcycle` in the `MASS` library. As a reminder, recreate the spline smoothing estimate using

```
library(MASS)
library(splines)
plot(mcycle)
fit1=smooth.spline(mcycle$times,mcycle$accel,df=10)
x.grid<-0:56
preds=predict(fit1,x.grid)
lines(x.grid,preds$y,lwd=2,col="black")
```

Use the code in `loess1.r` to obtain the local regression line at the point $x_0=28$ for the motorcycle data. You will need to complete some of the syntax. Once the code is working, try with `span=0.75` and again with $x_0=14$

The code in `loess2.r` evaluates the whole loess curve over a given x -grid to get a complete loess curve.

Use the R-command `loess()` to replicate your algorithm in in one step. Because you have fitted a *linear* local regression, you need to specify the argument `degree=1` in the `loess` function call. The

`loess()` function returns an list of class `loess`, which contains a vector called `$fitted`. These are the loess curve values evaluated at each of the data points. The `predict()` function is used to calculate the predicted values for any specified x -value(s).

2 Generalised additive Models

The Work data

Work through Lab 7.8.3 in James et al. starting on page 294, up to the command `plot(gam.lo.i)` on page 296.

The first example uses `lm()` to fit the model. Check that the function `gam()` with the same arguments outputs the same coefficients. Hints: if you have not already done so you need to install and start the `gam` package, and to obtain the coefficients of a statistical model use the function `coef()`

Important! The function `plot.gam` is now called `plot.Gam`

The College data

The `College` dataset is available in the `ISLR` package, the accompanying package to James et al. An introduction to the data set is given on page 55.

We will fit the variable `Accept`, the number of applications accepted, as the outcome variable. Obtain the mean and median and a histogram of `Accept`. What do you notice about the distribution of these data?

Fit a first model with `Private` (a yes/no factor variable) and `Apps` as a smoothing spline using 5 effective degrees of freedom.

Both variables are highly significant (`summary()`), but `Accept` and `Apps` have very skewed distributions. Repeat the GAM fitting with the logarithm of both of these variables.

Continue the analysis by adding in turn each of the following variables using spline smoothing, and using the `anova()` function to see if this variable gives a better model fit.

`F.Undergrad`, `Room.Board`, `Expend`, `PhD`, and `S.F.Ratio`.

Obtain the response plot for each of the fitted variables in your final including the partial residuals (`resid=TRUE`).

Use you final model to obtain a prediction for `Accept` at *Harvard University* and compare this with the observed value, NB remember that the logarithm of `accept` was fitted in the model.